



ComputerVision I

Image Processing

Roi Santos Mateos

Escola Técnica Superior de Enxeñaría, USC
Academic year 2025/26

MASTER IN ARTIFICIAL INTELLIGENCE



■ Image processing:

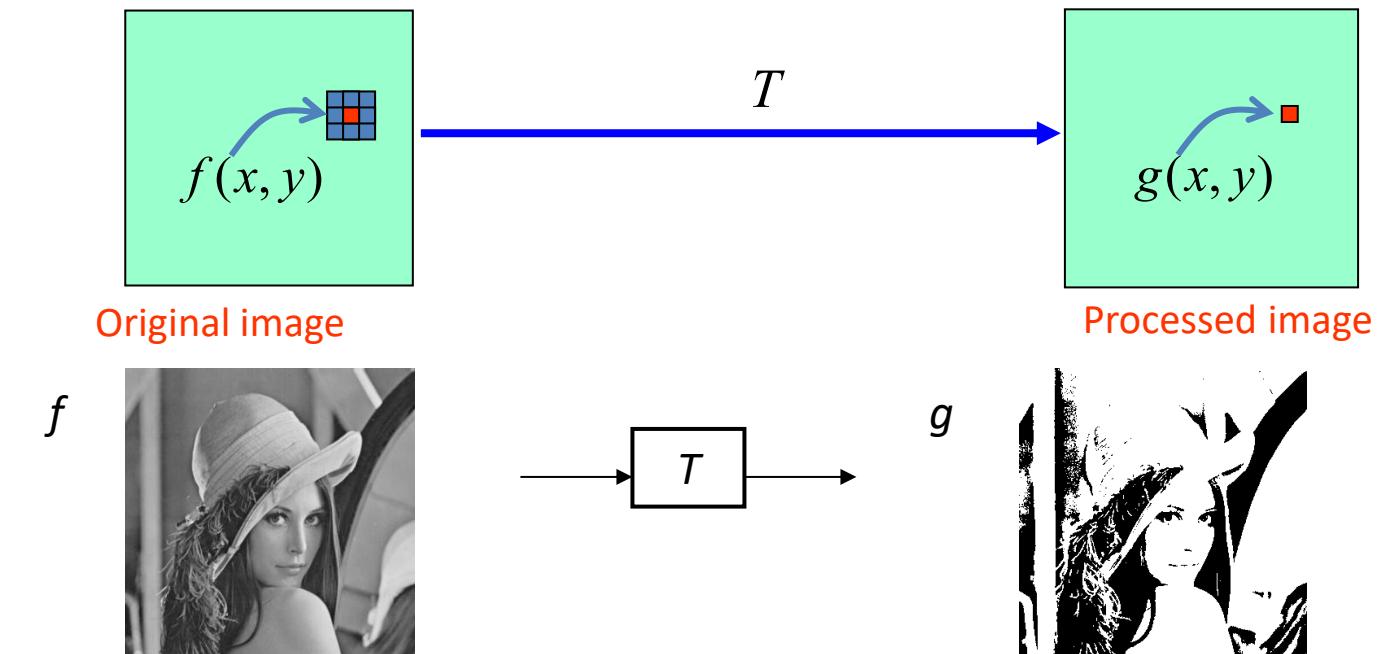
▷ It is a type of signal processing in which input is an image and output may be an image or a tensor of features associated with that image:

- Extracts useful information from images,
 - Features (edges, corners, blobs...)
- Improves image visual appearance
 - Enhancing, in-painting, de-noising
- Transforms image representation into another domain where image properties can be modified/enhanced easily
 - Discrete Fourier Transform, ...

Pixel-Level Image Processing

- Generate a new image taking into account only pixels values at the same location in the original image:

$$\xrightarrow{\text{Processed image}} g(x, y) = T[f(x, y)] \xleftarrow{\text{Original image}}$$



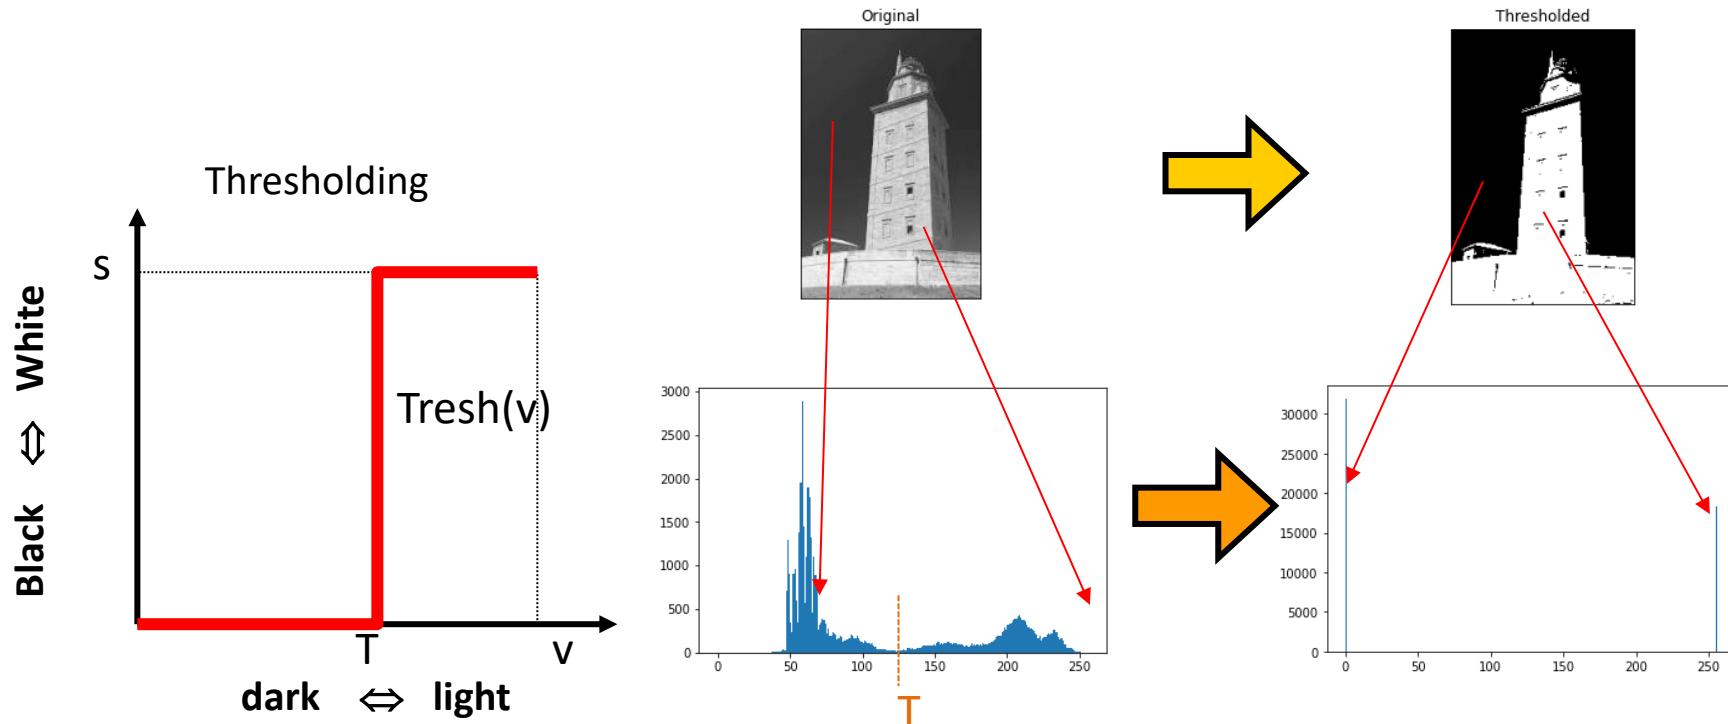
Histogram-Based Transformations

Thresholding

4

Thresholding:

- ▷ Separates foreground objects from background.
- ▷ Histogram ridges are related to big objects.
- ▷ Ridge height related to size.

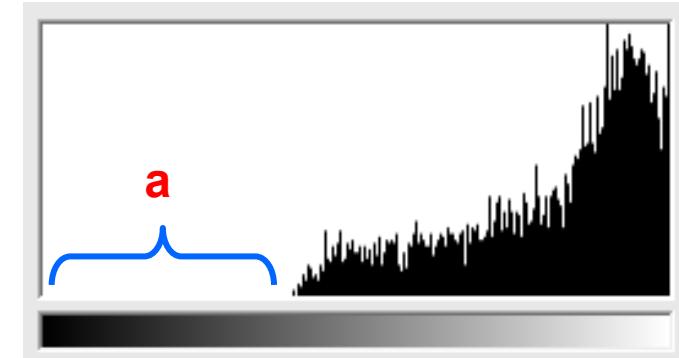
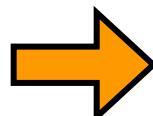
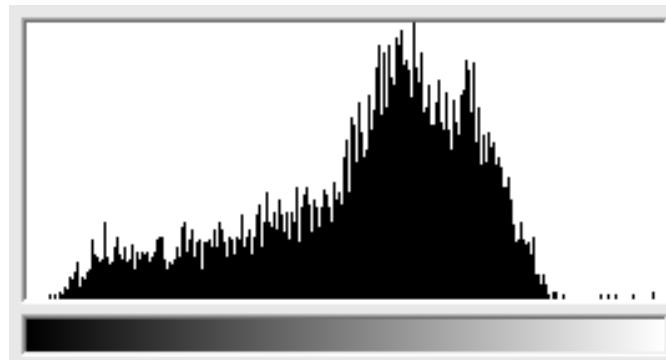
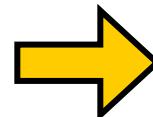


Histogram-Based Transformations

Arithmetic Transformation

5

- Adding a constant: $g(x, y) = f(x, y) + a$
 - ▷ Makes image brighter
 - ▷ Histogram “moves” to the right a pixels.



Histogram-Based Transformations

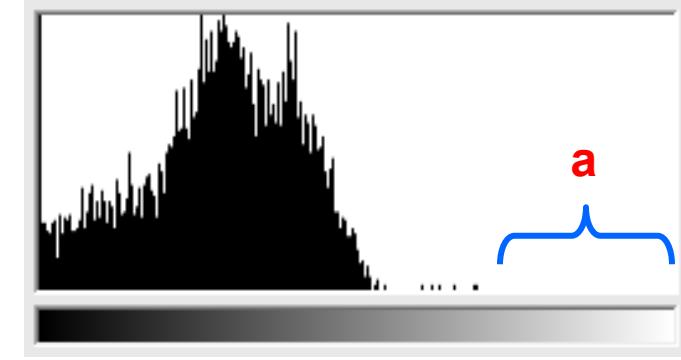
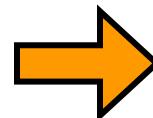
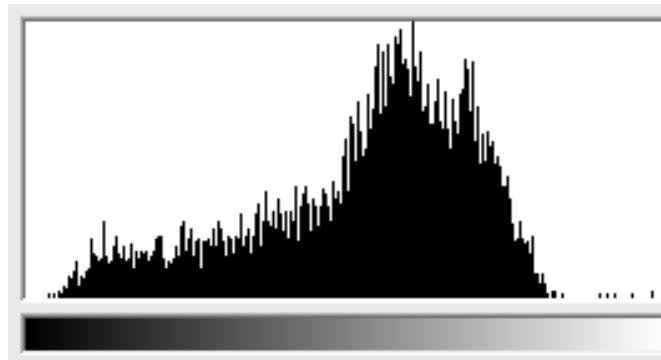
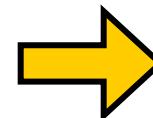
Arithmetic Transformation

6



- Subtracting a constant: $g(x, y) = f(x, y) - a$

- ▷ Makes image darker
- ▷ Histogram “moves” to the left a pixels.



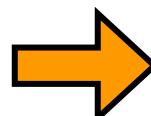
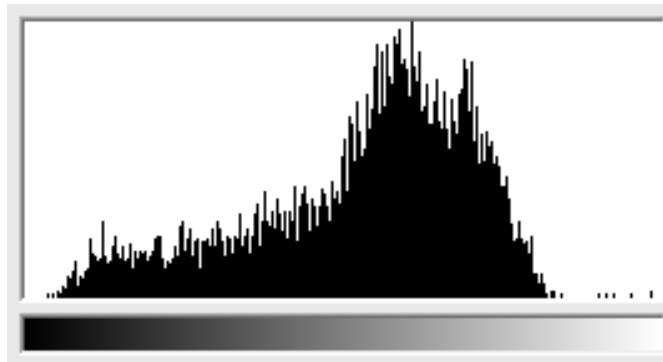
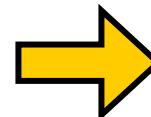
Histogram-Based Transformations

Arithmetic Transformation

7



- Multiplication by a constant: $g(x, y) = b \cdot f(x, y)$
 - ▷ Histogram “stretches” rightwards.



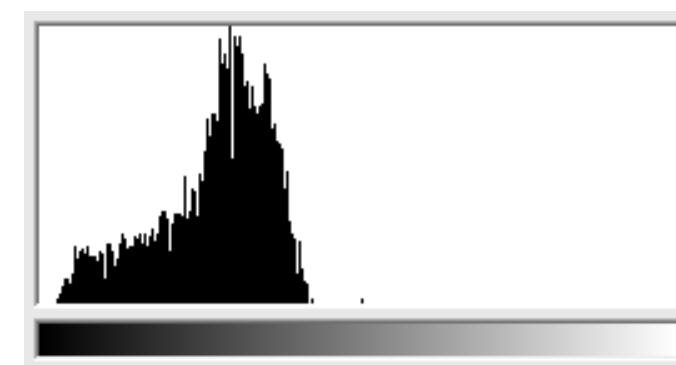
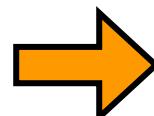
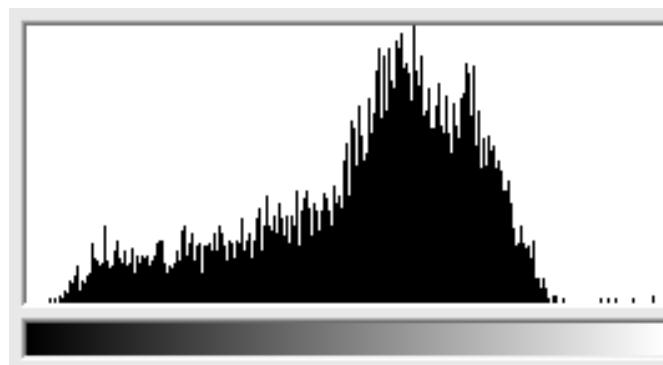
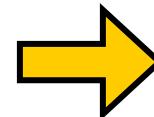
Histogram-Based Transformations

Arithmetic Transformation

8



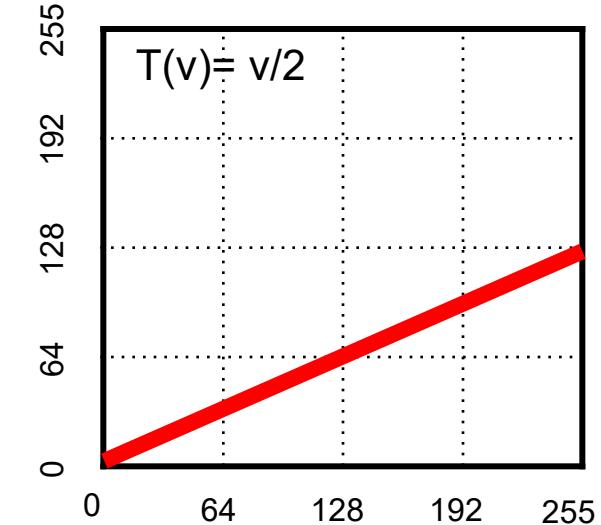
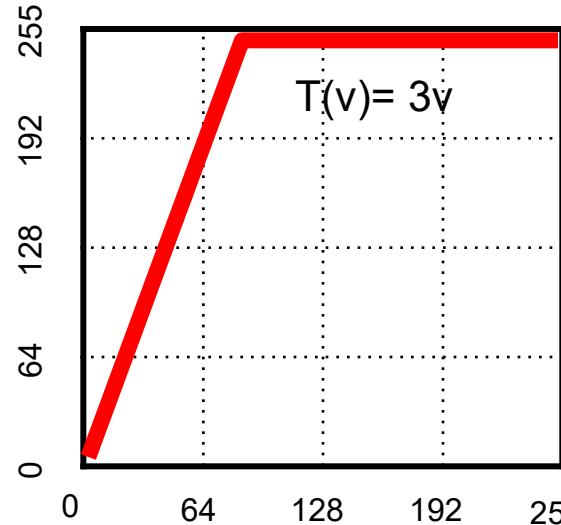
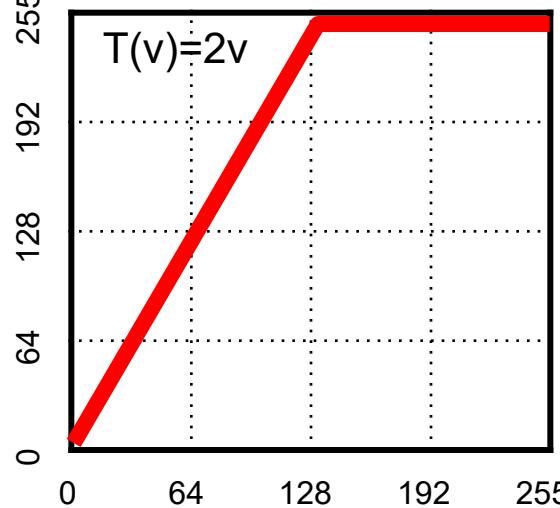
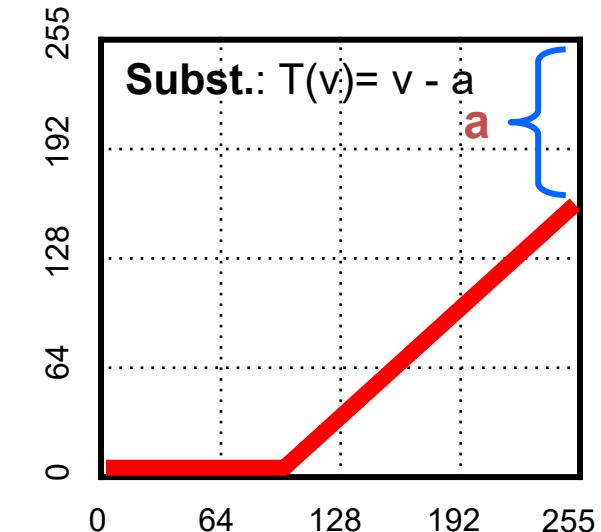
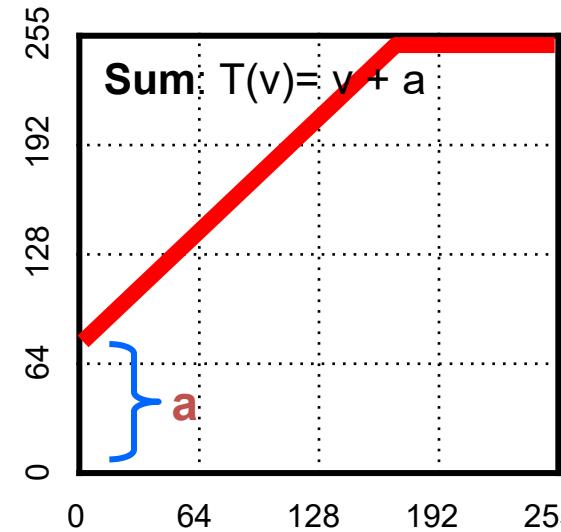
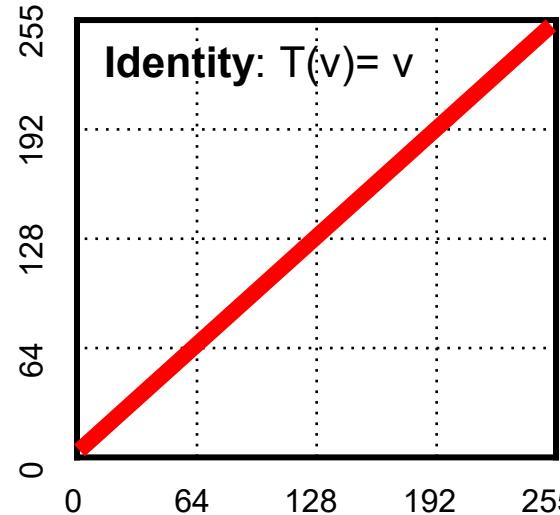
- Division by a constant: $g(x, y) = f(x, y) / b$
 - ▷ Histogram “shrinks”.



Histogram-Based Transformations

Histogram Transformations

9



Histogram-Based Transformations

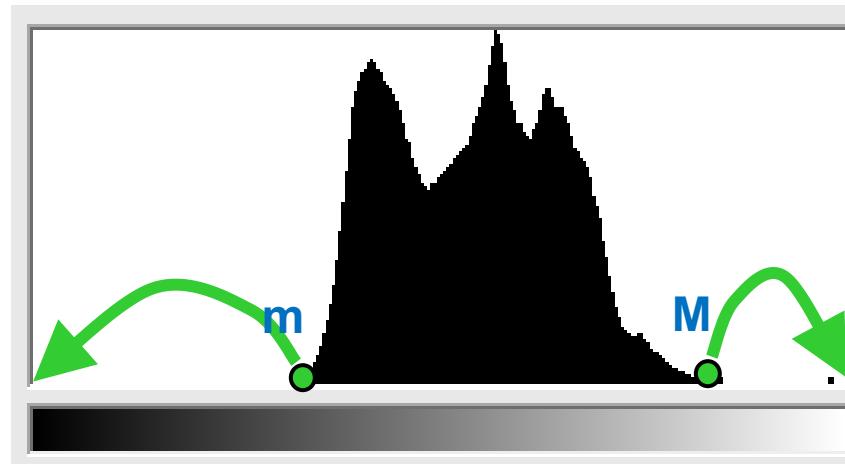
Image Enhancement

10



- Usually, it is interesting to stretch the histogram to enhance images
- Enhancement: defines a linear transformation of the histogram:
 - ▷ Finds the minimum gray-level value: **m**
 - ▷ Finds the maximum: **M**
 - ▷ **Performs an arithmetic transformation based on pixel intensity.**

$$T(v) = (v-m)*255/(M-m)$$



Histogram-Based Transformations

Image Enhancement

11



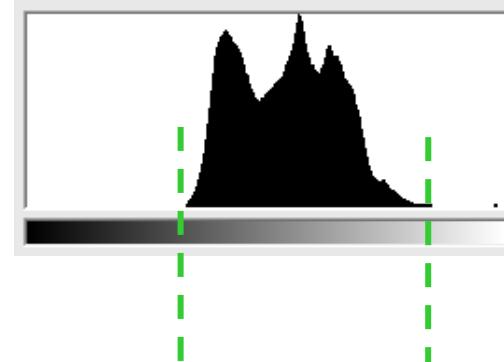
- Example: $m = 86$, $M = 214$



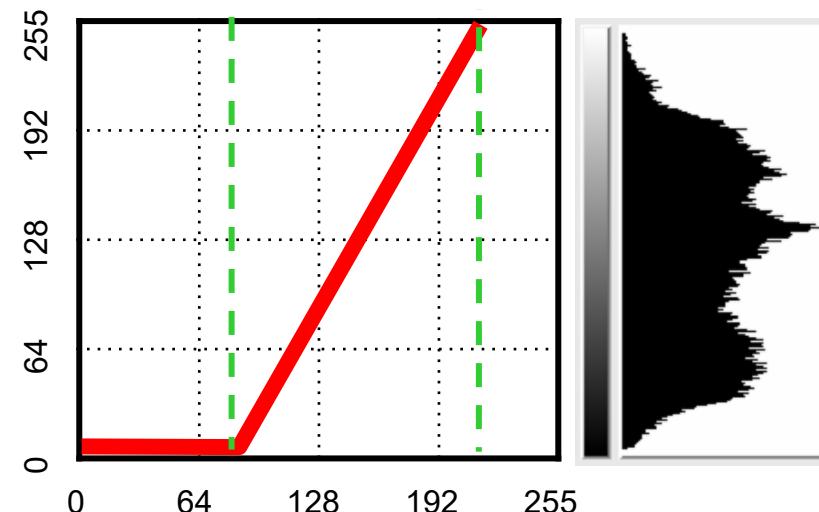
$$g(x,y) = (f(x,y)-86)*1.99$$



Histogram of f



$$T(v) = (v-86)*255/(214-86)$$



Histogram of g

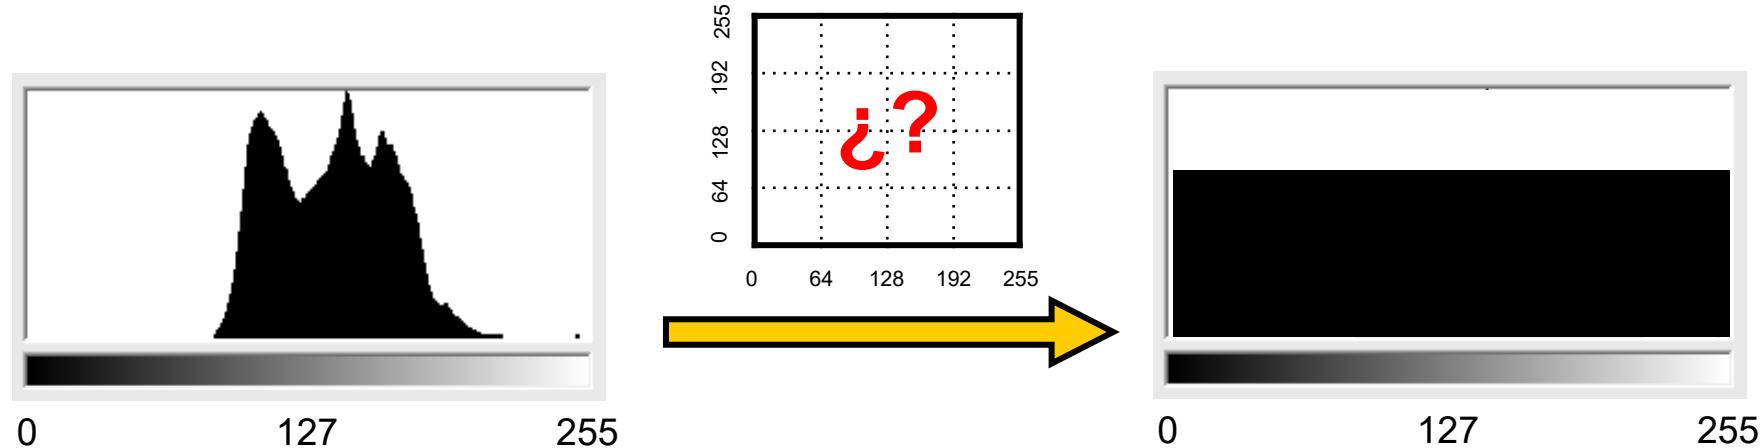
Histogram-Based Transformations

Equalization

12



- Equalization seeks to transform an image histogram in an (as much as possible) uniform one:
 - ▷ Goal: increasing contrast resolution in the output image



Histogram-Based Transformations

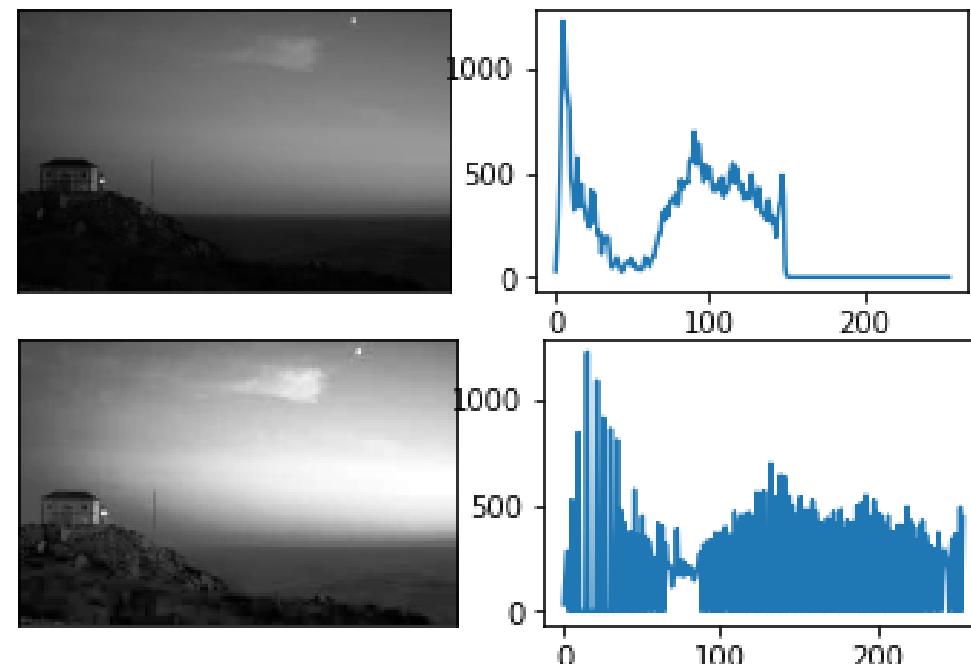
Equalization

13



■ Equalization steps:

- ▷ Compute histogram: gray-level Probability Distribution Function (PDF)
- ▷ Divide the gray-level range in **M intervals of the same probability** (same #pixels in all bins)
 - _ Intervals around ridges are narrower than intervals around valleys
- ▷ Split the histogram in M uniform bins and map original values (within each interval) in the new value in the corresponding bin (stretch or shrink)
- ▷ Transformed histogram will have **M intervals of gray-levels with same Cumulative Distribution Function (CDF)**



Histogram-Based Transformations

Equalization

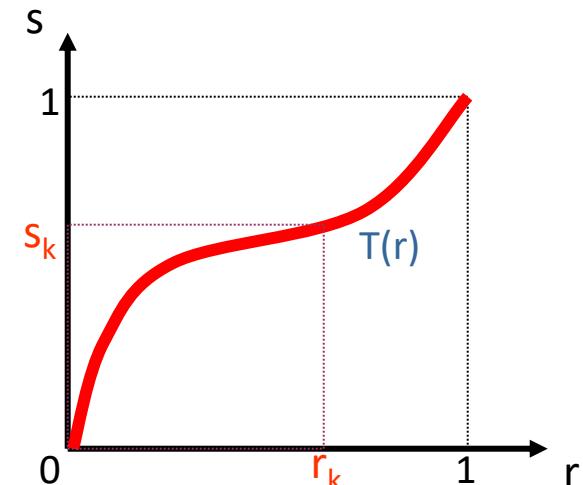
14



■ Fundamentals in continuous domain:

- ▷ Let r be a continuous variable that represents the intensity-levels of a normalized image in $[0,1]$.
- ▷ A transformation $s = T(r)$ is sought such that:
 1. It preserves intensity-level ordering :
 $T(r)$: monovalued and increasing monotone in $[0,1]$
 2. $0 \leq T(r) \leq 1, \forall 0 \leq r \leq 1$

This implies that stretching in some parts has to be done at the expenses of shrinking in others!



Histogram-Based Transformations

Equalization

15



- Transformation that fulfills conditions 1 and 2:

$$\mathbf{s} = \mathbf{T}(\mathbf{r}) = \int_0^{\mathbf{r}} \mathbf{p}_r(\mathbf{w}) d\mathbf{w}, \quad 0 \leq \mathbf{r} \leq 1 \text{ Cumulative Distribution Function (CDF)}$$

- In discrete domain:

$$s_k = T(r_k) = \sum_{j=0}^k h_r(r_j) = \sum_{j=0}^k \frac{n_j}{n} \quad \text{Cumulative Histogram}$$

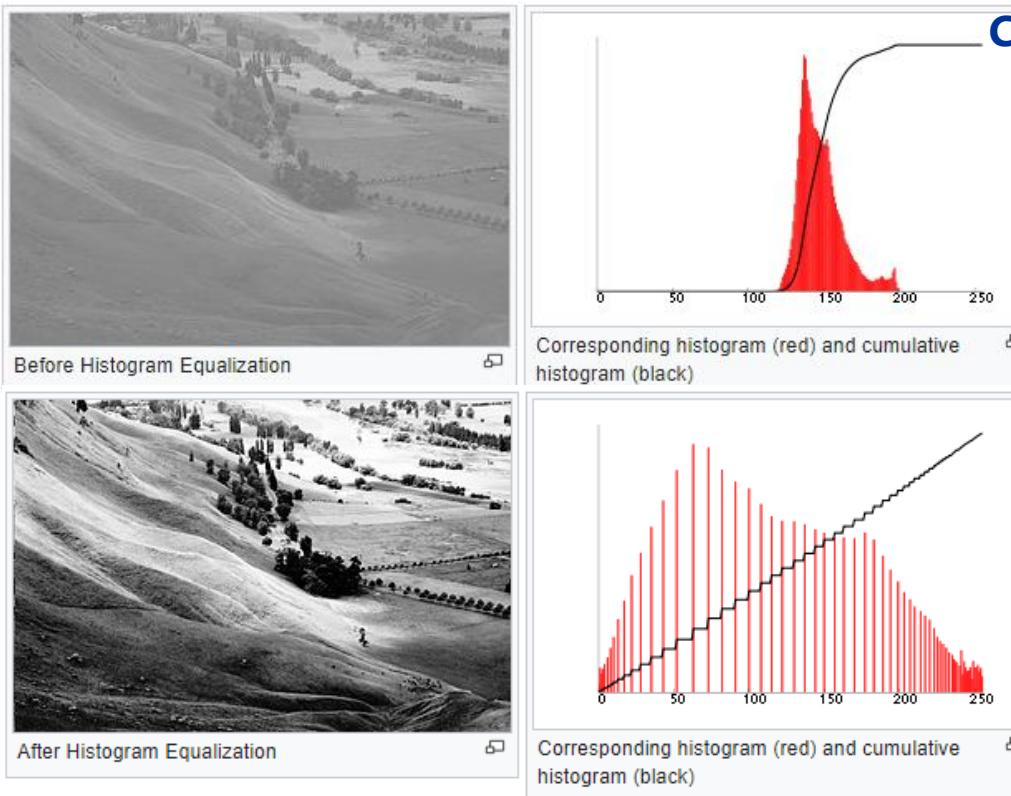
Histogram-Based Transformations

Equalization

16



- Example of discrete equalization:



Equalization redistributes gray-level bins such that all the intervals have similar number of pixels in their bins:

- 0.- M =Highest gray-scale value
- 1.- C = normalized cumulative histogram.
- 2.- $m_c = \min(C)$
- 3.- $s_r=T(r)= M*(c_r- m_c)/(1- m_c)$

To take advantage of all the gray scale range when $m_c > 0$

Point-Level Image Processing

Combination of images

17



■ Combination of images:

- ▷ Input: images A and B.
- ▷ Output: image R.

$$R(x, y) = f(A(x,y), B(x,y))$$

■ Multiple operators:

- ▷ Boolean: and, or, xor, not
- ▷ Arithmetic: sum, subtraction, multiplication
- ▷ Relational: max, min

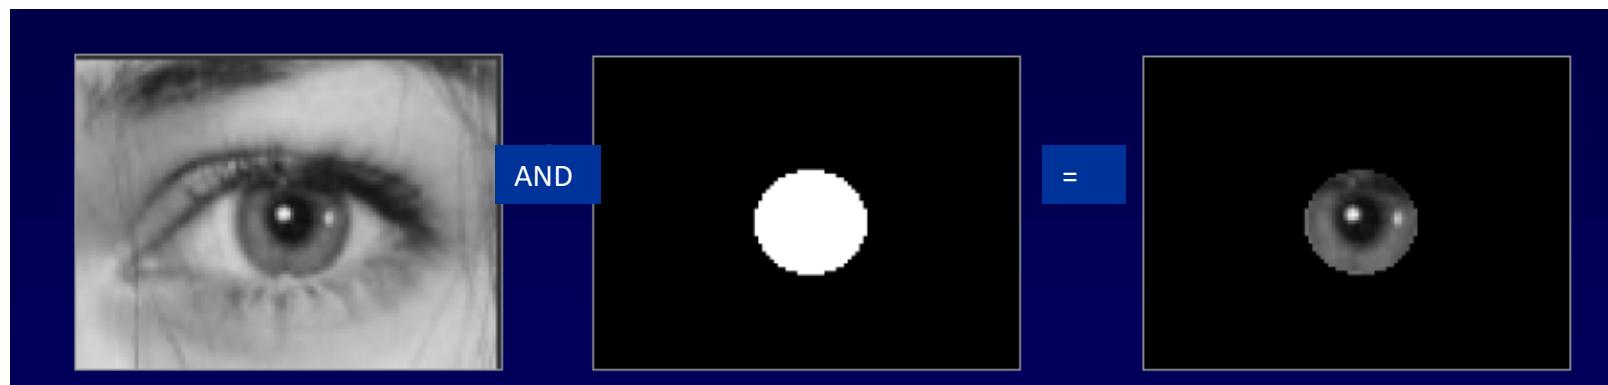
Boolean Operations



- Boolean bitwise operators:

- ▶ $R(x, y) = A(x,y) \text{ AND } B(x,y)$
- ▶ $R(x, y) = A(x,y) \text{ OR } B(x,y)$
- ▶ $R(x, y) = A(x,y) \text{ XOR } B(x,y)$
- ▶ $R(x, y) = \text{NOT } A(x,y) \text{ AND } B(x,y)$
- ▶ $R(x, y) = \text{NOT } A(x,y) \text{ OR } B(x,y) \dots$

- They are usually used when at least A is binary.



Arithmetic Operations



- Image subtraction: motion detection

$$\text{Abs}(\text{Image}_1 - \text{Image}_2) = \text{Result}$$
Three grayscale images illustrating motion detection. The first image shows a car approaching a railroad crossing. The second image shows the car at the crossing. The third image is a binary mask where the car is highlighted in white against a black background, indicating the detected motion.

- Mean: noise reduction:

$$\text{Mean}(\text{Image}_1, \text{Image}_2, \dots) = \text{Result}$$
Three images showing the mean operation. The first image is a noisy version of a wheel. The second image is a blurred version of the same wheel. The third image is a clean, sharp version of the wheel, representing the result of averaging the two images.

Local-Level Image Processing



- Capabilities of point operations are limited
- Local filtering: pixel transformation also depends on its neighborhood
 - ▷ Dependency is expressed by means of a mask

$w(-1,-1)$	$w(-1,0)$	$w(-1,1)$
$w(0,-1)$	$w(0,0)$	$w(0,1)$
$w(1,-1)$	$w(1,0)$	$w(1,1)$

Mask coefficients showing coordinate arrangement, around transformed pixel coordinates (0,0)

Local-Level Image Processing

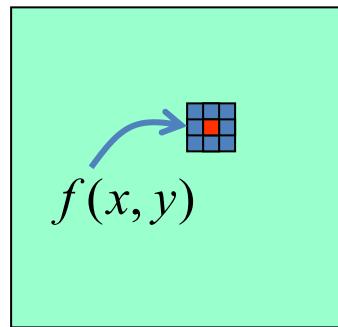
2D Convolution

21



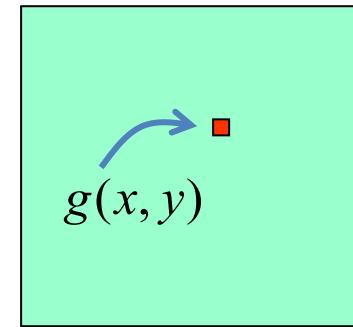
Convolution:

$$g(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b \omega(s, t) f(x-s, y-t)$$



Input Image, f

$$g = \omega * f$$



Output Image, g

Local-Level Image Processing

2D Convolution

22



- Local-Level Image Processing
 - ▷ Convolution

2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

Input Image, f

*

1	-1	-1
1	2	-1
1	1	1

Convolution kernel, ω

=

5	4	4	-2
9	6	14	5
11	7	6	5
9	12	8	5

Output Image, g

Equivalently: kernel rotation + kernel sliding
+ summation of the multiplication of
overlapped image and kernel values.

- ▷ First step: kernel rotation

1	-1	-1
1	2	-1
1	1	1

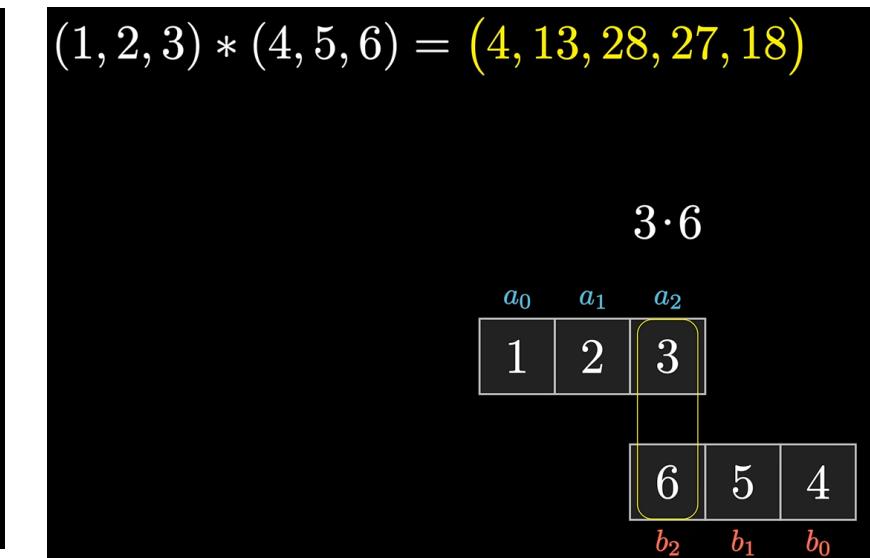
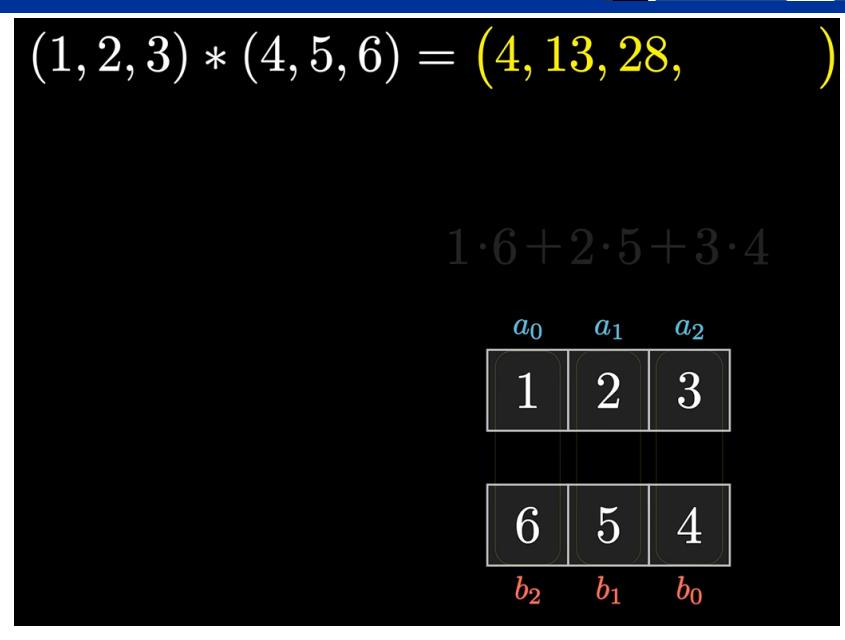
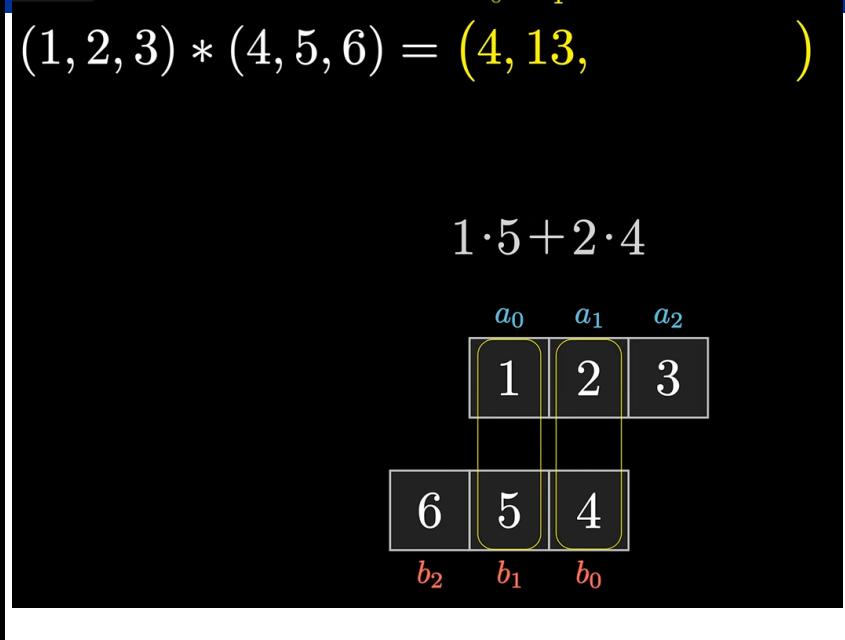
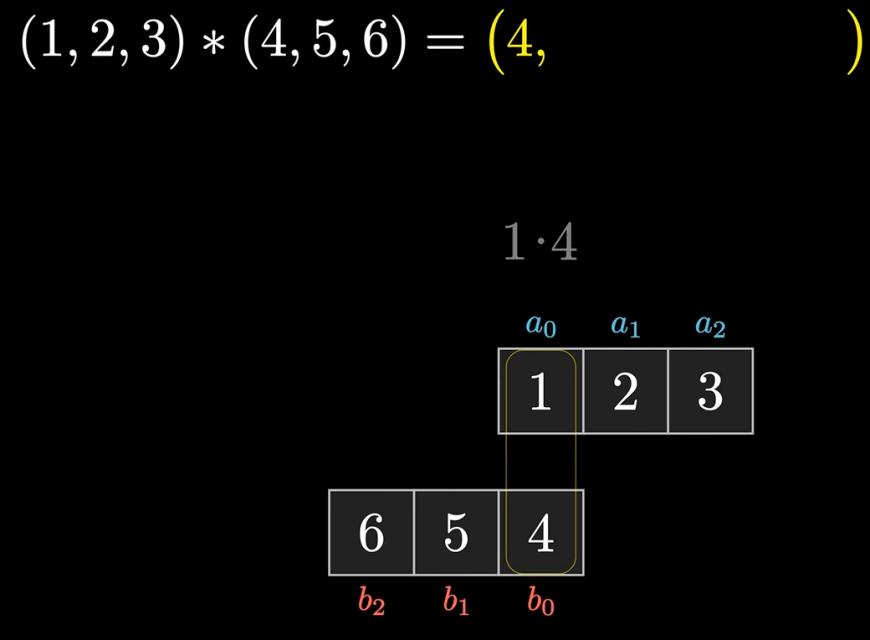
Rotate 180°

1	1	1
-1	2	1
-1	-1	1

Local-Level Image Processing

2D Convolution

23



Local-Level Image Processing

2D Convolution

24



$$(1, 2, 3) * (4, 5, 6) = (4, 13, 28, 27, 18)$$

$c_0 \quad c_1 \quad c_2 \quad c_3 \quad c_4$

```
(base) grant [ ~/cs/videos ]$ ipython
Python 3.9.7 | packaged by conda-forge | (default, Sep 29 2021, 19:24:02)
Type 'copyright', 'credits' or 'license' for more information
IPython 7.29.0 -- An enhanced Interactive Python. Type '?' for help.
```

```
In [1]: import numpy as np
```

```
In [2]: np.convolve([1, 2, 3], [4, 5, 6])
Out[2]: array([ 4, 13, 28, 27, 18])
```

```
In [3]:
```

a_0	a_1	a_2
1	2	3

b_2	b_1	b_0
6	5	4

Local-Level Image Processing

2D Convolution

25



- Second step:
 - kernel sliding
 - summation of the multiplications of overlapped image and kernel values

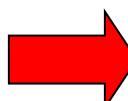
Convolution kernel, ω

1	1	1
-1	2	1
-1	-1	1

Input Image, f

2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

1	1	1		
-1	4	2	2	3
-1	-2	1	3	3
2	2	1	2	
1	3	2	2	



Output Image, g

5			

Local-Level Image Processing

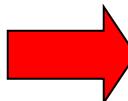
2D Convolution



- Next step:
 - kernel sliding
 - summation of the multiplications of overlapped image and kernel values

1	1	1
-1	2	1
-1	-1	1

1	1	1	
-2	4	2	3
-2	-1	3	3
2	2	1	2
1	3	2	2



2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

5	4		

Output Image, g

Local-Level Image Processing

2D Convolution

27

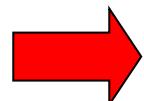


- ▷ Next step:
 - kernel sliding
 - summation of the multiplications of overlapped image and kernel values

1	1	1
-1	2	1
-1	-1	1

2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

2	2	2	3
-2	2	3	3
-2	-2	1	2
1	3	2	2



5	4	4	-2
9	6		

Local-Level Image Processing

2D Convolution

28



- Final output image g:

5	4	4	-2
9	6	14	5
11	7	6	5
9	12	8	5

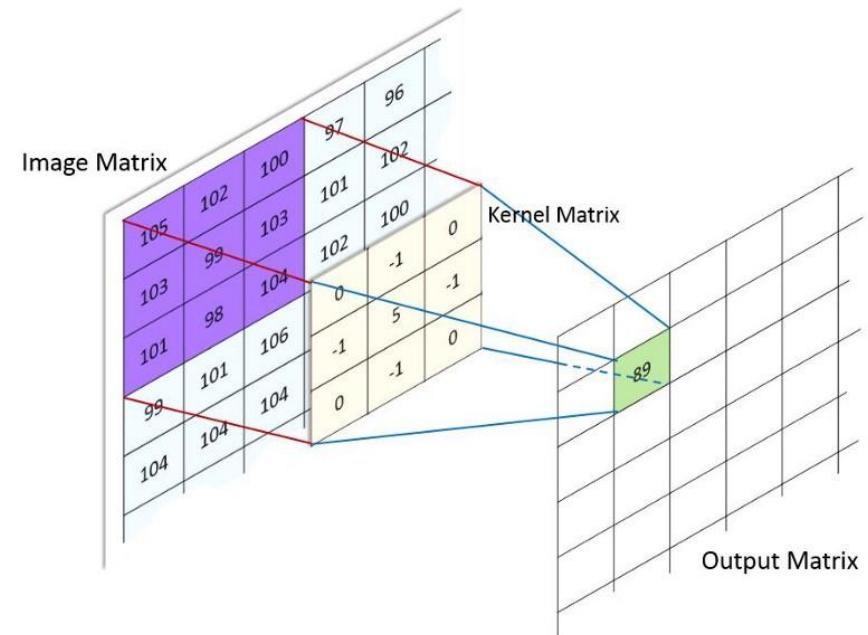
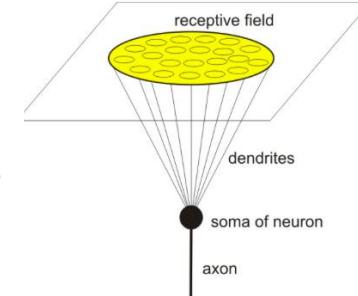
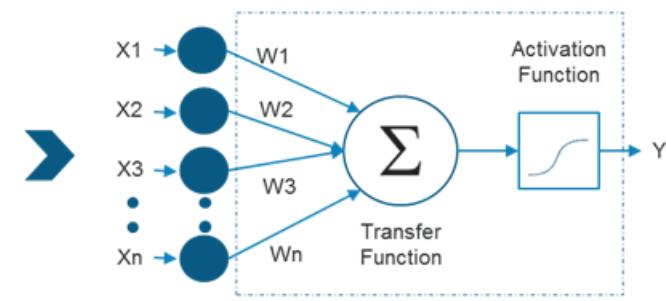
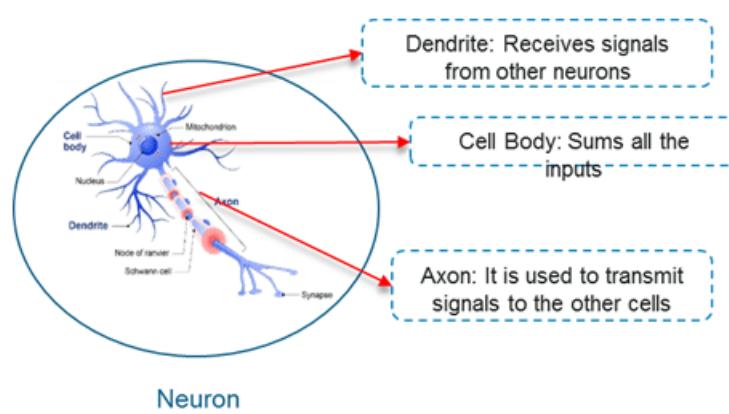
Local-Level Image Processing

2D Convolution

29



Convolution: from biological Inspiration to implementation



Local-Level Image Processing

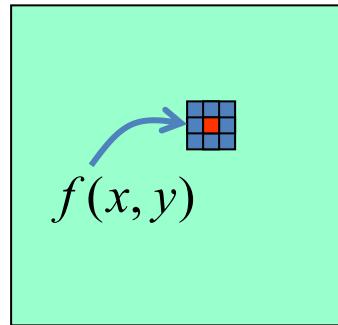
2D Correlation

30



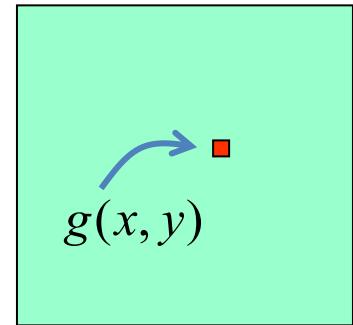
Correlation

$$g(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b \omega(s, t) f(x + s, y + t)$$



Input Image f

$$g = \omega \circ f$$



Output Image g

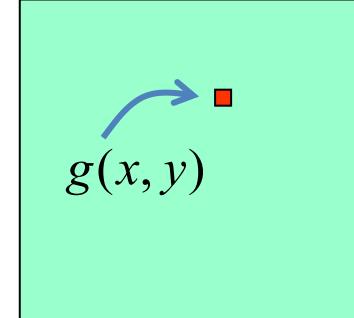
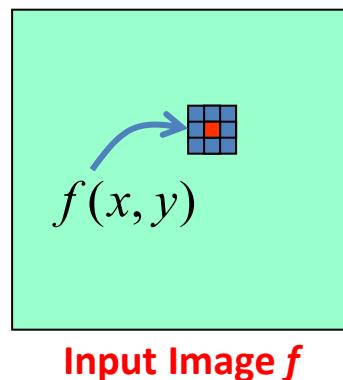
Local-Level Image Processing

2D Correlation

31



Local-Level Image Processing



Output Image g

Convolution kernel ω

1	-1	-1
1	2	-1
1	1	1

Don't Rotate

Equivalently: kernel sliding + summation of the multiplication of overlapped image and kernel values.

Input Image f

2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

Local-Level Image Processing

2D Correlation

32



1	-1	-1
1	2	-1
1	1	1

1	-1	-1		
1	4	-2	2	3
1	2	1	3	3
2	2	1	1	2
1	3	2	2	

Input Image, f

2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

output Image, g

5			

Local-Level Image Processing

2D Correlation

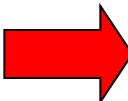
33



1	-1	-1
1	2	-1
1	1	1

1	-1	-1	
2	4	-2	3
2	1	3	3
2	2	1	2
1	3	2	2

Input Image f



2	2	2	3
2	1	3	3
2	2	1	2
1	3	2	2

Output Image g

5	10		

Local-Level Image Processing

2D Correlation

34



Local-Level Image Processing

$$\begin{matrix} 1 & -1 & -1 \\ 1 & 2 & -1 \\ 1 & 1 & 1 \end{matrix} \quad o \quad \text{Kernel } \omega$$

$$\begin{matrix} 2 & 2 & 2 & 3 \\ 2 & 1 & 3 & 3 \\ 2 & 2 & 1 & 2 \\ 1 & 3 & 2 & 2 \end{matrix} \quad \text{Input Image } f$$

$$= \begin{matrix} 5 & 10 & 10 & 15 \\ 3 & 4 & 6 & 11 \\ 7 & 11 & 4 & 9 \\ -5 & 4 & 4 & 5 \end{matrix} \quad \text{Final output Image } g$$



Correlation vs Convolution

- They are simple operations, but extremely useful
 - ▷ Linear
 - ▷ Shift invariant
- They do the same job when kernels are symmetrical
- Key difference: convolution is associative and commutative:
 - ▷ That is, if F and G are filters, then
 - ▷ $F^*(G^*I) = (F^*G)^*I$
- In general:
 - ▷ Convolution is used for image processing operations such as smoothing, edge detection
 - ▷ Correlation is used to match a template to an image.
- Correlation: $g(x, y) = \sum_{s,t} \omega(s, t) f(x + sy + t)$
- Convolution: $g(x, y) = \sum_{s,t} \omega(-s - t) f(x + sy + t) \leftarrow \text{kernel is flipped.}$
- So $g_{conv} = f * \omega = f \circ \tilde{\omega}$ with $\tilde{\omega}(s, t) = \omega(-s, -t)$.



■ Convolution ($f * g$)

- ▷ **Associative:** $(f * g) * h = f * (g * h)$
- ▷ **Commutative:** $f * g = g * f \leftarrow$ (mathematically, for infinite-support signals)
- ▷ **Distributive:** $f * (g + h) = f * g + f * h$
- ▷ Identity: $f * \delta = f$

■ Cross-correlation ($f \circ g$)

- ▷ **Not commutative** in general: $f \circ g \neq g \circ f$
- ▷ **Not associative** in general
- ▷ Relation to convolution: $f \circ g = f * \tilde{g}$, where $\tilde{g}(x, y) = g(-x - y)$ and complex-conjugate if needed.

■ Special case: if the kernel is **real and symmetric** ($g = \tilde{g}$,then **correlation = convolution**, so commutativity holds.

Convolution and Correlation



- Image filtering: compute function of local neighborhood at each position
- Really important!
 - ▷ Enhance images
 - Denoise, resize, increase contrast, etc.
 - ▷ Extract information from images
 - Texture, edges, distinctive points, etc.
 - ▷ Detect patterns
 - Template matching
 - ▷ Deep Convolutional Networks

Convolution

- Example: box filter

 - ▷ Replaces each pixel with an average of its neighborhood

 - Achieve smoothing effect (remove sharp features)

$$w = \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

f

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

g

	0	10	20	30	30	30	20	10		
	0	20	40	60	60	60	40	20		
	0	30	60	90	90	90	60	30		
	0	30	50	80	80	90	60	30		
	0	30	50	80	80	90	60	30		
	0	20	30	50	50	60	40	20		
	10	20	30	30	30	30	20	10		
	10	10	10	0	0	0	0	0		

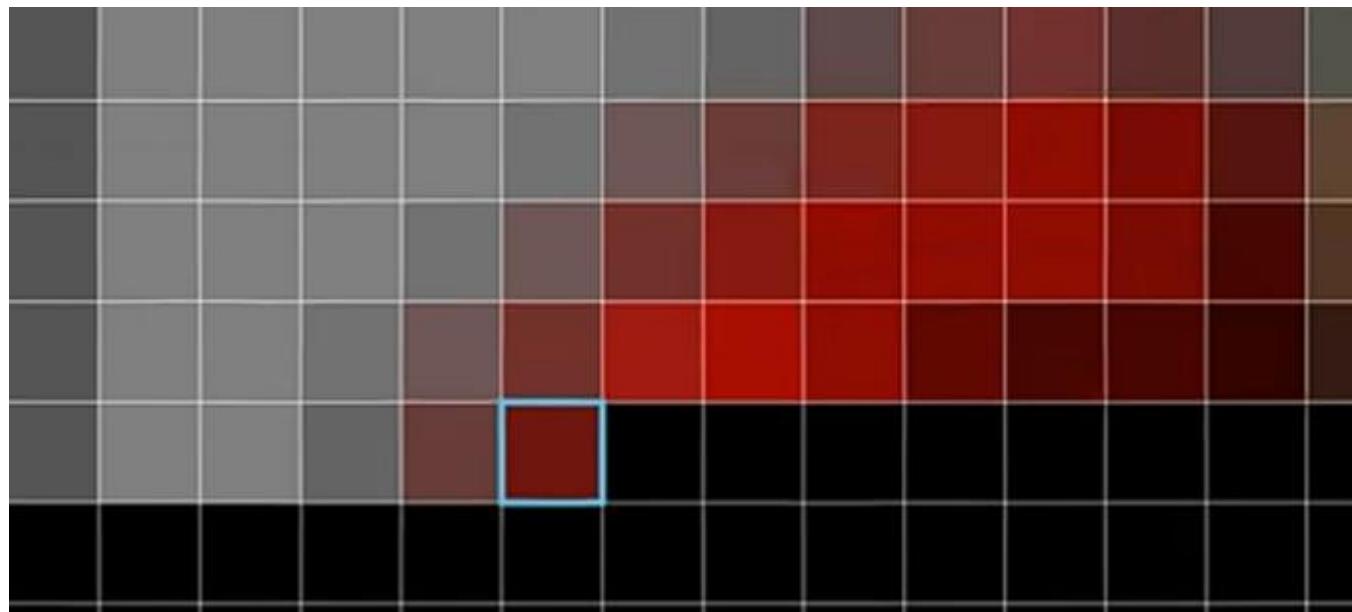
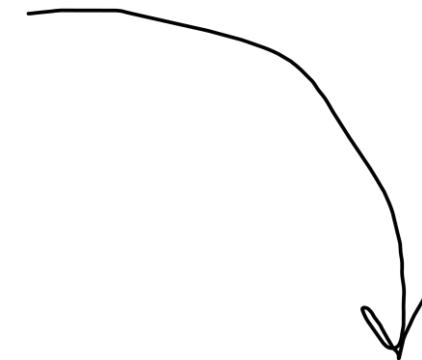
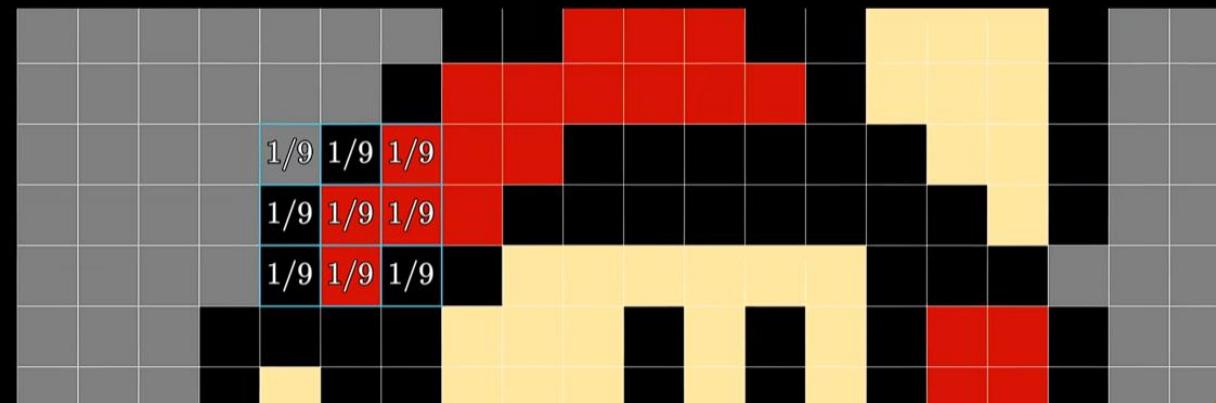
$$g[m, n] = \sum_{k,l} w[k, l] f[m - k, n - l]$$

Convolution

39



$$\frac{1}{9} \begin{bmatrix} 0.5 \\ 0.5 \\ 0.5 \end{bmatrix} + \frac{1}{9} \begin{bmatrix} 0.0 \\ 0.0 \\ 0.0 \end{bmatrix} + \frac{1}{9} \begin{bmatrix} 0.8 \\ 0.0 \\ 0.0 \end{bmatrix} + \frac{1}{9} \begin{bmatrix} 0.0 \\ 0.0 \\ 0.0 \end{bmatrix} + \frac{1}{9} \begin{bmatrix} 0.8 \\ 0.0 \\ 0.0 \end{bmatrix} + \frac{1}{9} \begin{bmatrix} 0.8 \\ 0.0 \\ 0.0 \end{bmatrix} + \frac{1}{9} \begin{bmatrix} 0.0 \\ 0.0 \\ 0.0 \end{bmatrix} + \frac{1}{9} \begin{bmatrix} 0.8 \\ 0.0 \\ 0.0 \end{bmatrix} + \frac{1}{9} \begin{bmatrix} 0.0 \\ 0.0 \\ 0.0 \end{bmatrix}$$

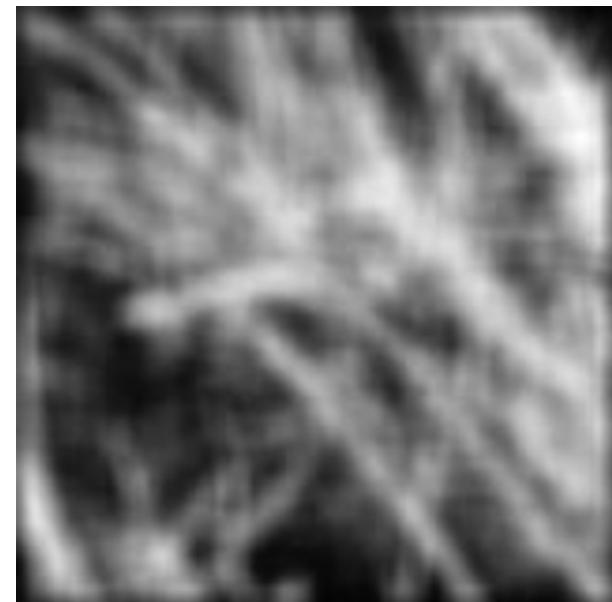


Convolution

40



- Example: box filter
 - ▷ Replaces each pixel with an average of its neighborhood
 - Achieve smoothing effect (remove sharp features)

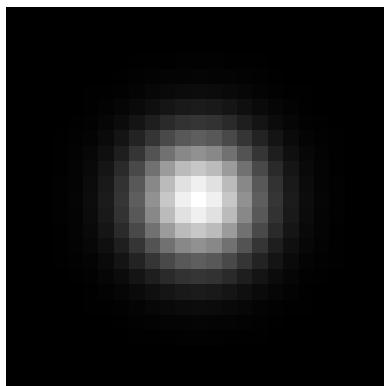
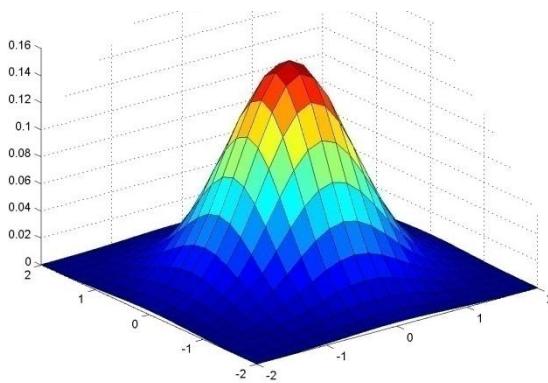


Convolution

41



- Example: Gaussian filter
 - ▷ Weight contributions of neighboring pixels by nearness
 - Achieve smoothing effect (remove sharp features)



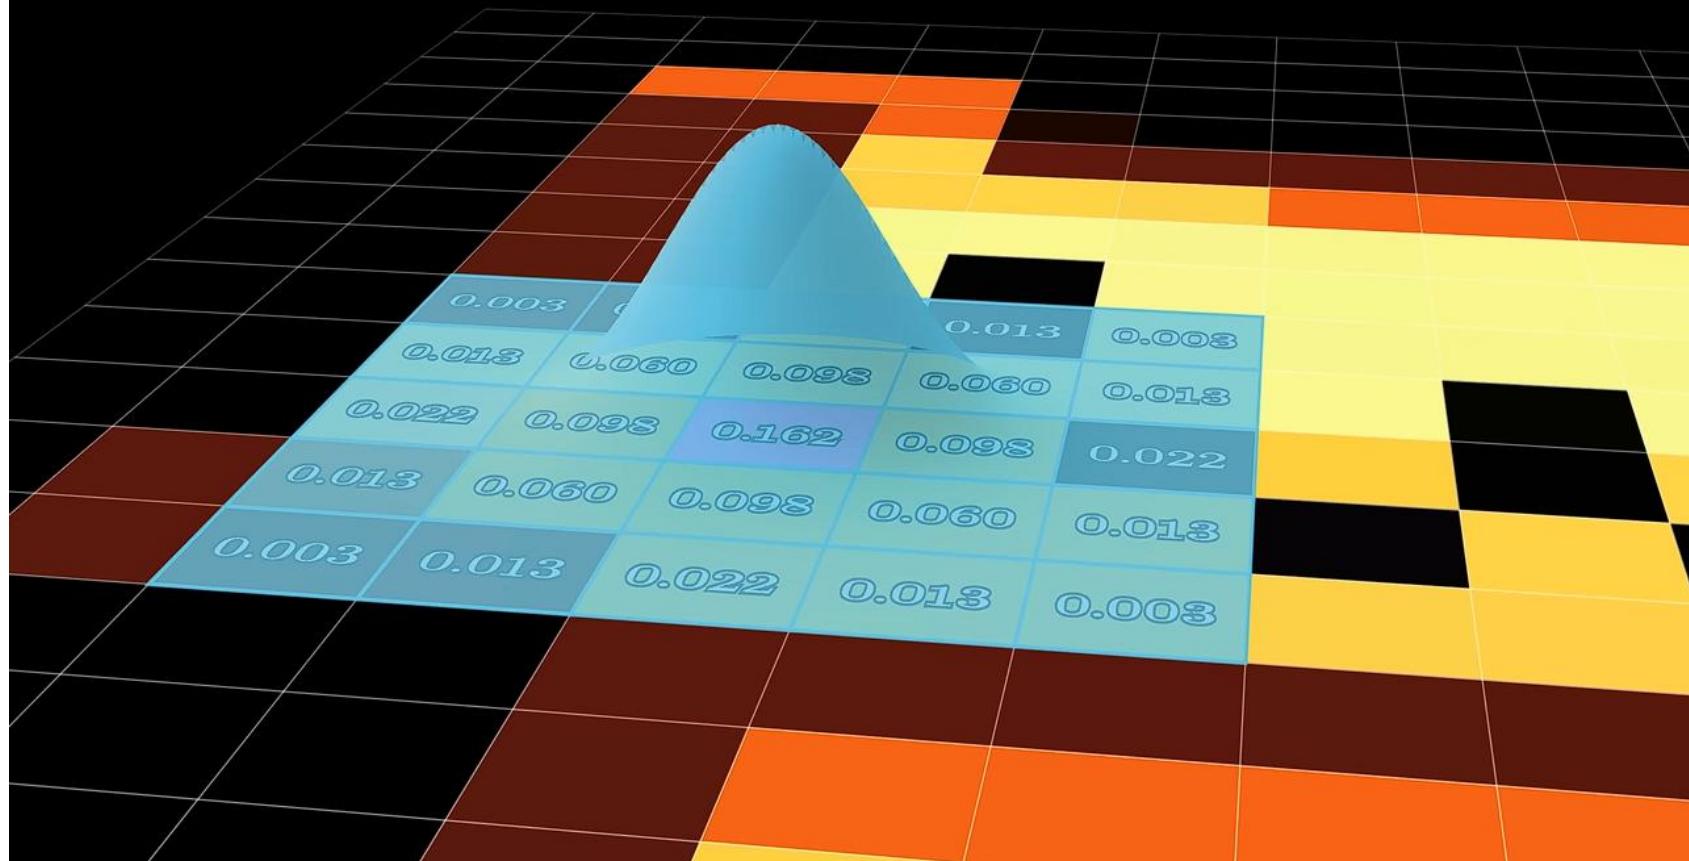
0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

$5 \times 5, \sigma = 1$

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$



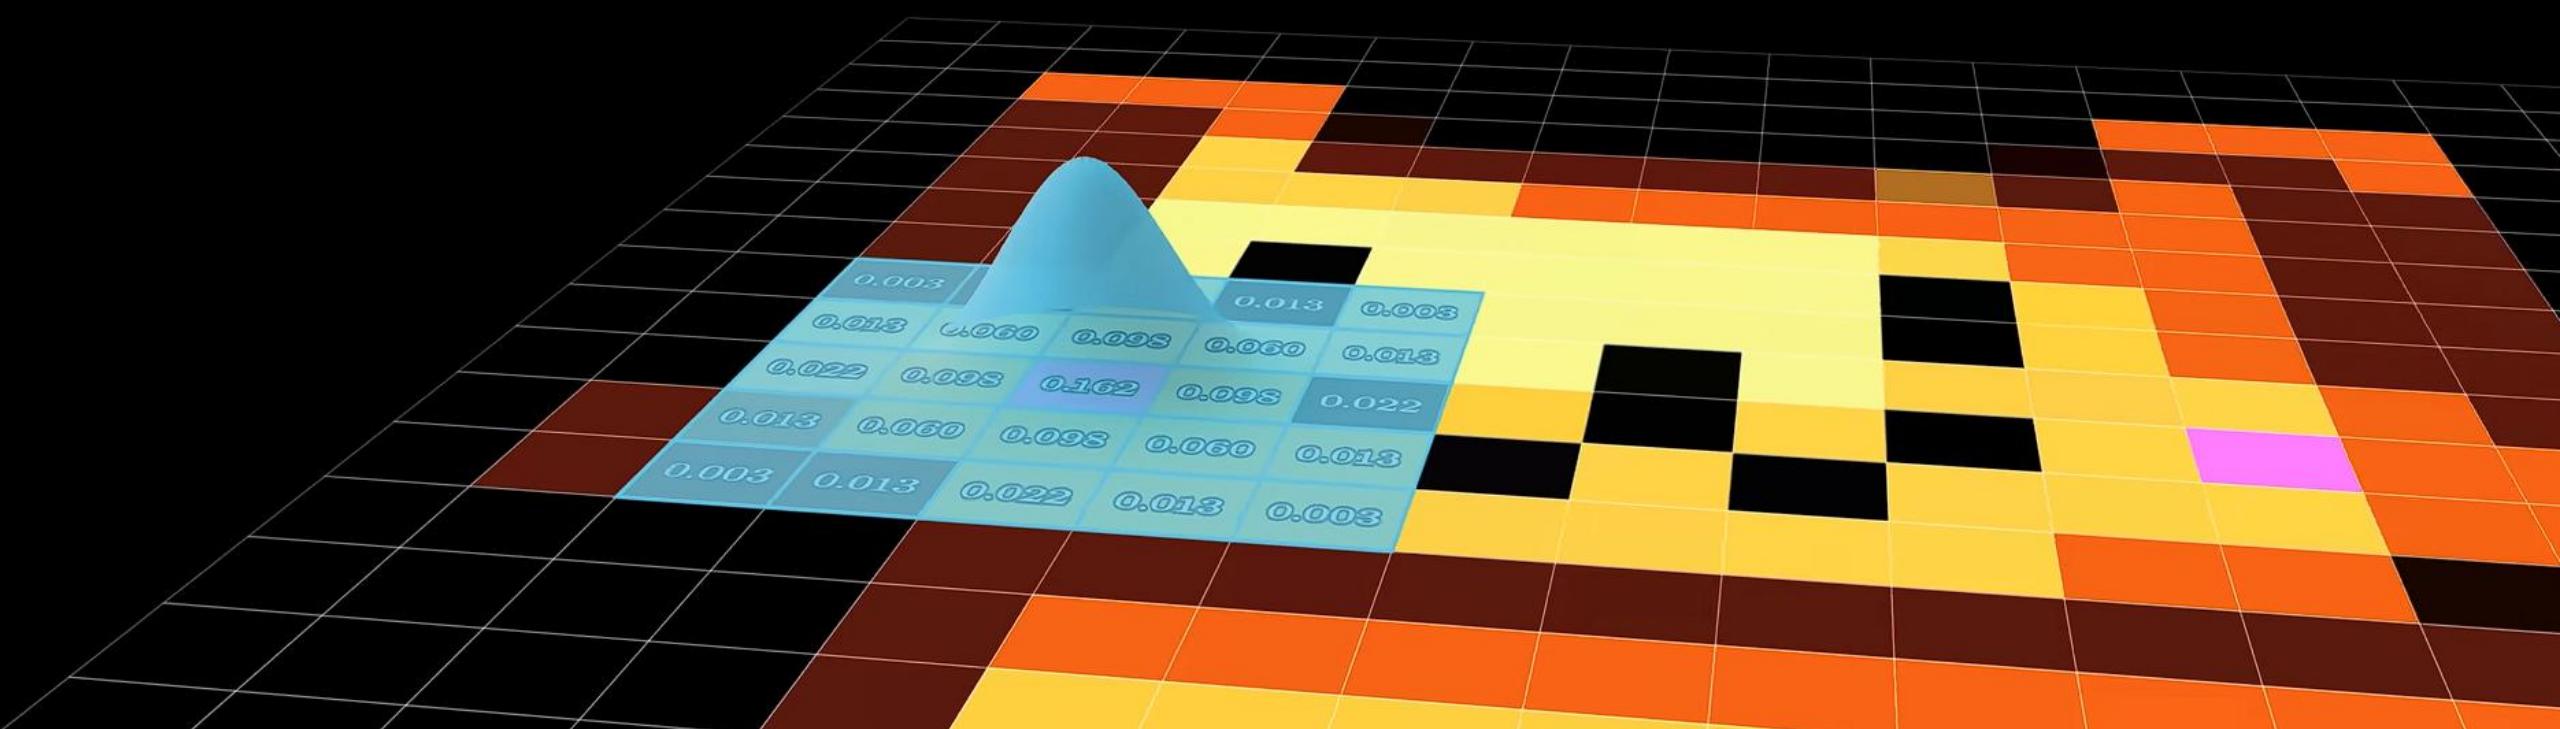
Gaussian distribution



Convolution



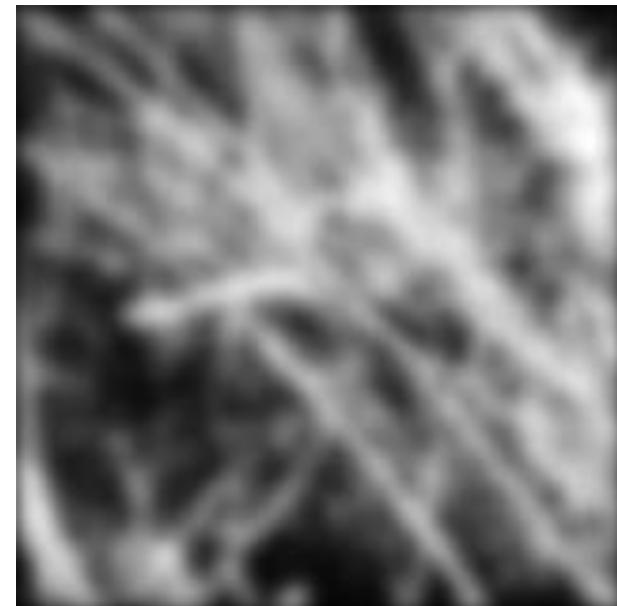
0.003 + 0.013 + 0.022 + 0.013 + 0.003 + 0.013 + 0.060 + 0.098 + 0.060 +
0.013 + 0.022 + 0.098 + 0.162 + 0.098 + 0.022 + 0.013 + 0.060 + 0.098 +
0.060 + 0.013 + 0.003 + 0.013 + 0.022 + 0.013 + 0.003



Convolution



- Example: Gaussian filter
 - ▷ Weight contributions of neighboring pixels by nearness
 - Achieve smoothing effect (remove sharp features)

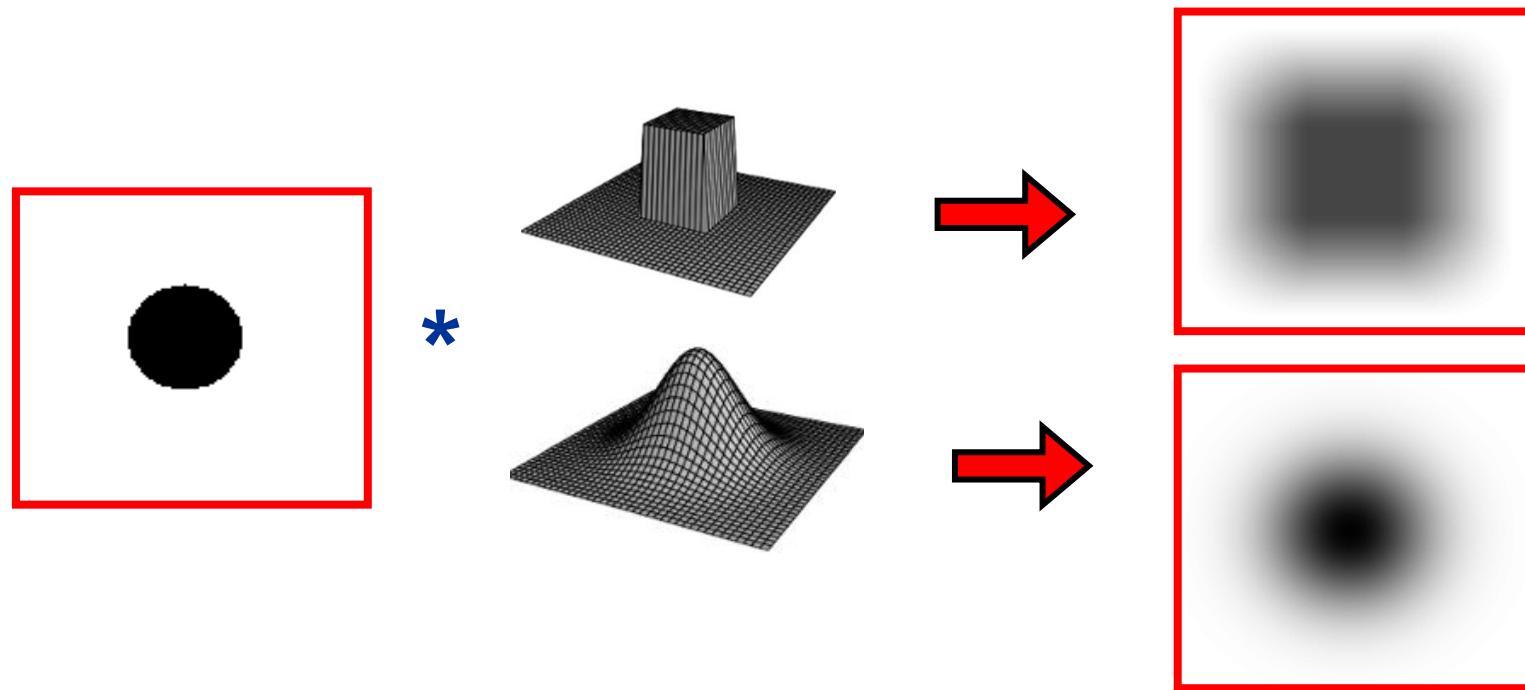


Denoising



■ Gaussian average (GA) vs mean filter (MF):

- ▷ GA is isotropic
- ▷ MF smooths further along diagonals than along rows and columns.
- ▷ GA weights decay gradually to zero
- ▷ MF weights have an abrupt cut-off which leaves discontinuities in the smoothed image.



Gaussian Kernel

46

- Standard deviation σ , parametrizes the smoothing level.
 - ▷ Greater values: wider bell, larger kernel size, higher smoothing.
 - ▷ Smaller values: narrower bell, lower smoothing.

- Approximation to discrete Gaussian (kernel 1D):
 - ▷ Pascal triangle.

$$\begin{matrix} & & & & & 1 \\ & & & & & 1 & 1 \\ & & & & & 1 & 2 & 1 \\ & & & & & 1 & 3 & 3 & 1 \\ & & & & & 1 & 4 & 6 & 4 & 1 \\ & & & & & 1 & 5 & 10 & 10 & 5 & 1 \end{matrix}$$

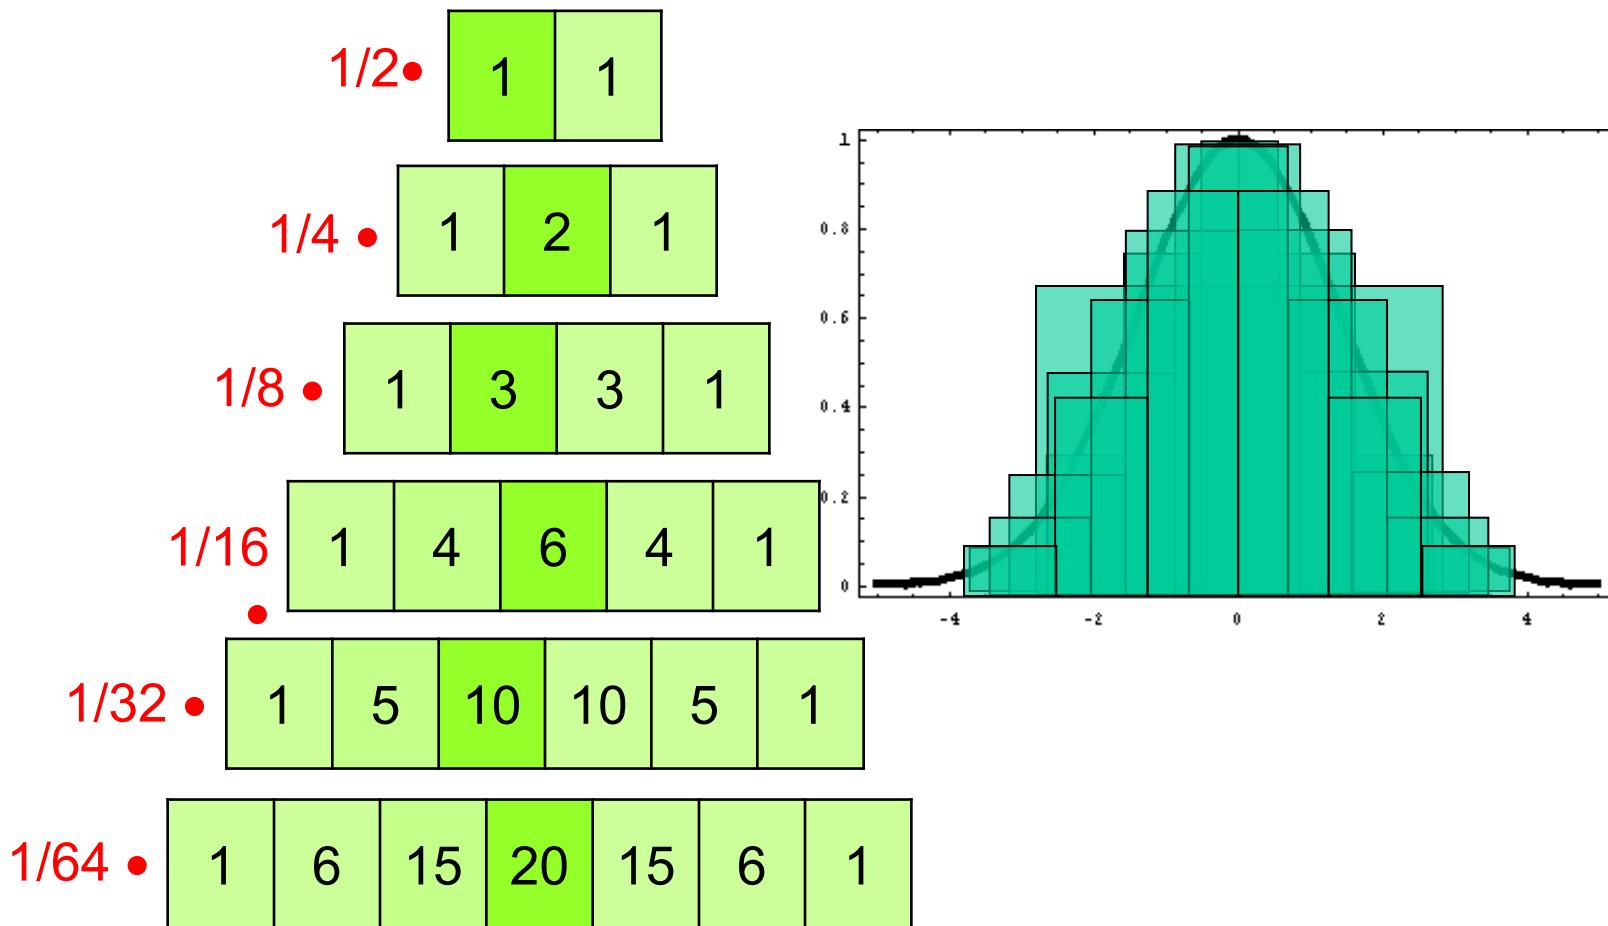
nf •

nf: normalization factor

Gaussian kernel



- Pascal triangle rows correspond to different discrete versions of a Gaussian



Gaussian Filters



- Remove “high-frequency” components from the image (low-pass filter)
- Convolution with itself is another Gaussian
 - ▷ Double smoothing with small-width kernel, give same result as smoothing with larger-width kernel
 - ▷ Convolving two times with Gaussian kernel of width σ is same as convolving once with kernel of width $\sigma\sqrt{2}$
- *Separable* kernel: Factors into product of two 1D Gaussians:

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

- ▷ Discrete domain: The filter factors into a product of 1D filters:

$$\begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array} = \begin{array}{|c|} \hline 1 \\ \hline 2 \\ \hline 1 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline \end{array}$$

Correlation

Matching with templates

49

- Goal: find  in image
- Correlation of an image with the window template

$$R(x, y) = \sum_{x', y'} (T(x', y') \cdot I(x + x', y + y'))$$



I



T



R

(Probable matchings can be obtained by thresholding)

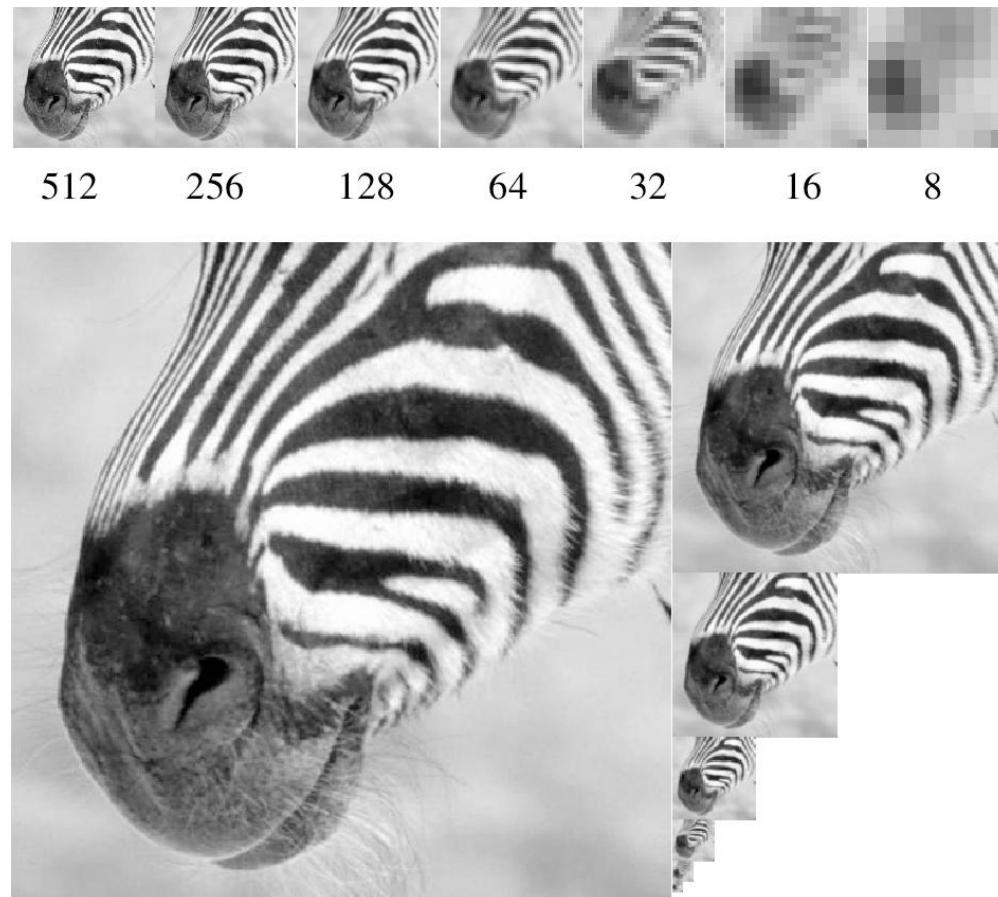


Location of maximum

Image Pyramid



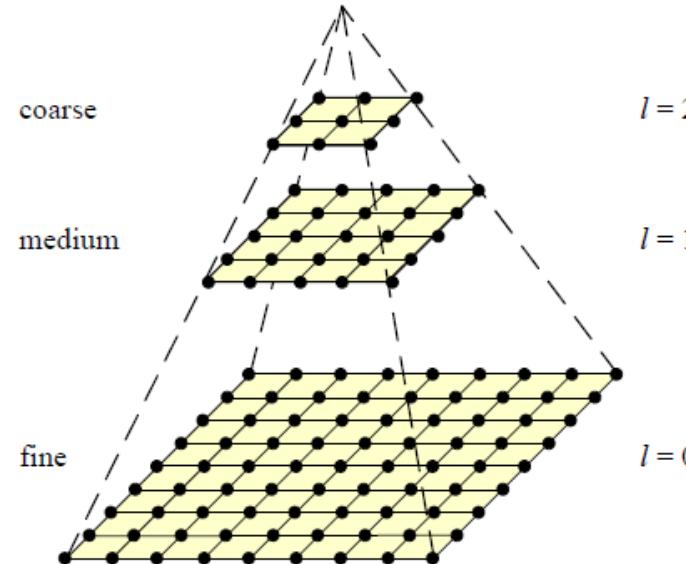
- What if we want to find larger or smaller objects/parts?



Template Matching with Image Pyramids



- What if we want to find larger or smaller objects/parts?
- Input: Image, Template
 - 1. Match template at current scale
 - 2. Downscale image
 - 3. Repeat 1-2 until desired minimum scale (size of interest)
 - 4. Take responses above some threshold, perhaps with non-maxima suppression



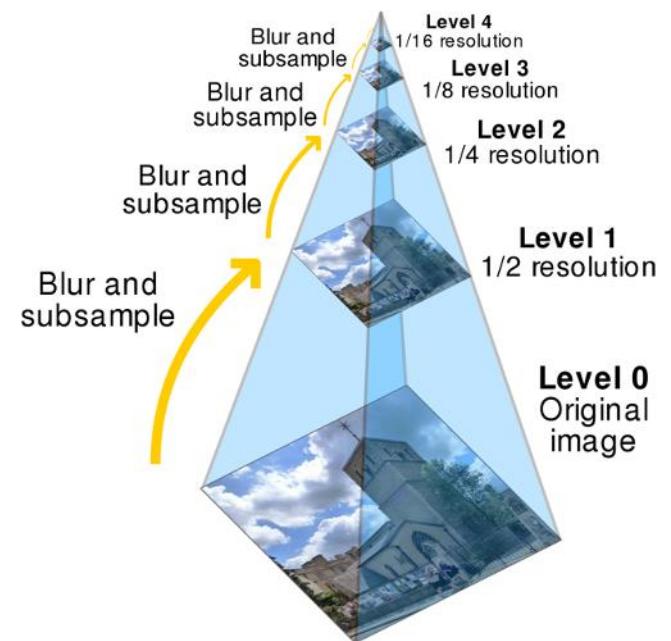
Template Matching with Image Pyramids



■ Image Pyramids

- ▷ Level 0: original image
- ▷ For $i=1$ to L
 - Gaussian smoothing Level $i-1$
 - Level i : Remove odd cols and rows

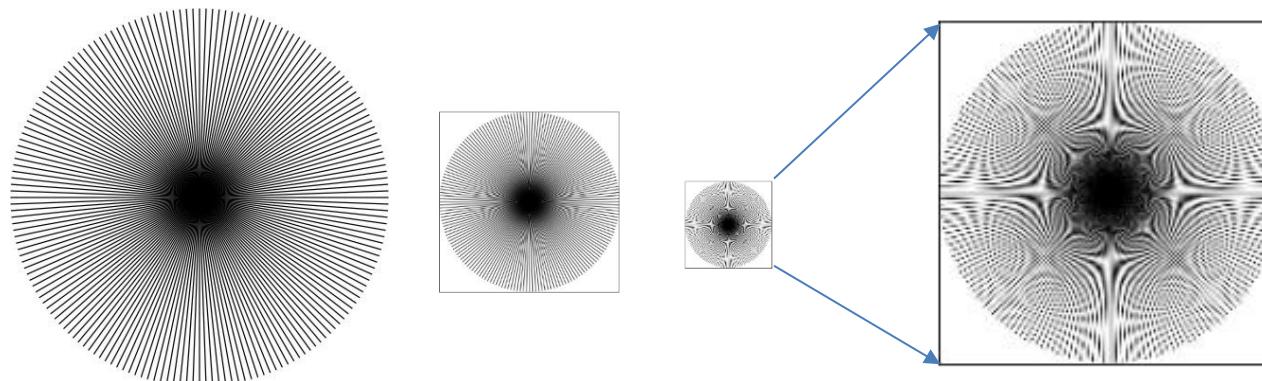
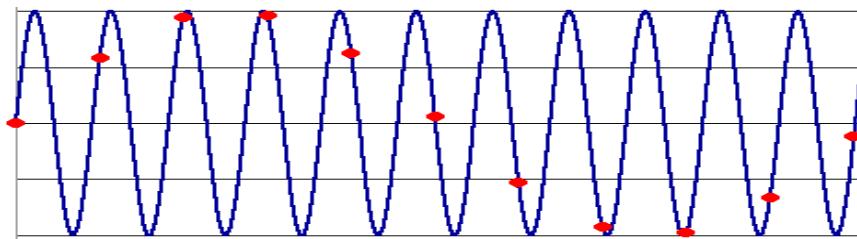
■ Why is smoothing necessary?



Template Matching with Image Pyramids



- Why is smoothing necessary?
 - ▷ Because of the aliasing effect
 - Occurs when sampling rate is not high enough to capture the amount of detail in an image
 - Can give the wrong signal/image—an *aliasing*



Template Matching with Image Pyramids

54



- Why is smoothing necessary?
 - ▷ To avoid aliasing:
 - Sampling rate $\geq 2 * \text{max frequency in the image}$
 - said another way: \geq two samples per cycle
 - Solution: smooth the image, *then* subsample

